

A Neural Network Approach for Real-Time High-Dimensional Optimal Control

Mines AMS Colloquium

Nov 12, 2021

Derek Onken
derekonken.com



EMORY UNIVERSITY

FOUNDED 1836

Collaborators and Acknowledgments



Xingjian Li
Emory



Lars Ruthotto
Emory



Samy Wu Fung
Emory/UCLA/Mines



Levon Nurbekyan
UCLA



Stan Osher
UCLA

Funding:



UNITEDHEALTH GROUP®

Special thanks: Organizers and staff of IPAM Long Program MLP 2019 and NVIDIA.

Overview

- **Background**

- ▶ Problem
- ▶ Pontryagin Maximum Principle (PMP)
- ▶ Hamilton–Jacobi–Bellman Partial Differential Equation (HJB)

- **Mathematical Formulation**

- ▶ Shock-Robustness
- ▶ HJB Penalizers

- **Neural Networks (NNs)**

- ▶ Model Formulation
- ▶ Numerics

- **Results**

- ▶ 150-Dimensional Swarm Trajectory Planning
- ▶ Quadcopter with Complicated Dynamics

- **Conclusion**

Optimal Control (OC) Problem

Corridor Problem

Consider two *centrally-controlled* agents that navigate through a corridor/valley between two hills to fixed targets

Assume

- We have control over the agents' velocities (the *control*)


Want

- Shortest paths, e.g. the geodesics (*optimality*)
- No collisions
- Agents to reach targets at final time

Multi-Agent Formulation

Consider n agents initially at $x_1, \dots, x_n \in \mathbb{R}^q \Rightarrow \mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^d$

Agents follow trajectories $z_{\mathbf{x}}(t)$ during time $t \in [0, T]$



Initial

$$z_{\mathbf{x}}(0) = \mathbf{x} = \begin{bmatrix} -2 \\ -2 \\ 2 \\ -2 \end{bmatrix} \left\{ \begin{array}{l} \text{agent 1} \\ \text{agent 2} \end{array} \right.$$

Target

$$\mathbf{y} = \begin{bmatrix} 2 \\ 2 \\ -2 \\ 2 \end{bmatrix}$$

Terminal Cost

$$G(z_{\mathbf{x}}(T)) = \frac{\alpha_1}{2} \|z_{\mathbf{x}}(T) - \mathbf{y}\|^2$$

for multiplier $\alpha_1 \in \mathbb{R}$

Trajectories Governed by Differential Equation

The state z_x depends on the control u_x and previous state via the system

$$\begin{aligned}\partial_t z_x(t) &= f(t, z_x(t), u_x(t)), \quad z_x(0) = x \\ \text{For Corridor:} \quad &= u_x(t) \quad (\text{the velocity})\end{aligned}\tag{1}$$

where

- time $t \in [0, T]$
- initial state $x \in \mathbb{R}^d$
- admissible controls $U \subset \mathbb{R}^a$
- $f: [0, T] \times \mathbb{R}^d \times U \rightarrow \mathbb{R}^d$ models the evolution of the state $z_x: [0, T] \rightarrow \mathbb{R}^d$ in response to the control $u_x: [0, T] \rightarrow U$

Running Cost

Running costs where z_i and u_i are the state and control for the i th agent, respectively

$$\begin{aligned} L(t, z(t), u(t)) &= E(z(t), u(t)) + \alpha_2 Q(z(t), u(t)) + \alpha_3 W(z(t), u(t)) \\ &= \sum_{i=1}^n \underbrace{E_i(z_i(t), u_i(t))}_{\text{For Corridor: } \frac{1}{2}\|u_i(t)\|^2} + \alpha_2 \sum_{i=1}^n \underbrace{Q_i(z_i(t), u_i(t))}_{\text{sum of Gaussians}} + \alpha_3 \sum_{j \neq i} \underbrace{W_{ij}(z_i(t), z_j(t))}_{\text{piecewise Gaussian repulsion}} \end{aligned}$$

for multipliers $\alpha_2, \alpha_3 \in \mathbb{R}$ and

- E_i is the energy of an agent,
- Q_i represents any obstacles or terrain,
- W_{ij} are the interaction costs between homogeneous agents i and j with radius r

$$W_{ij}(z_i, z_j) = \begin{cases} \exp\left(-\frac{\|z_i - z_j\|_2^2}{2r^2}\right), & \|z_i - z_j\|_2 < 2r \\ 0, & \text{otherwise} \end{cases}$$

Optimal Control (OC) Problem

$$\begin{aligned} \text{Running Cost: } L(s, \cdot) &= E(\cdot) + \alpha_2 Q(\cdot) + \alpha_3 W(\cdot) \\ \text{Terminal Cost: } G(z_x(T)) &= \frac{\alpha_1}{2} \|z_x(T) - \mathbf{y}\|^2 \end{aligned}$$

Goal: Find the control that incurs minimal cost¹

$$\Phi(t, \mathbf{x}) = \inf_{\mathbf{u}_x} \left\{ \int_t^T L(s, z_x(s), \mathbf{u}_x(s)) ds + G(z_x(T)) \right\} \quad (2)$$

- $\Phi(t, \mathbf{x}) \in \mathbb{R}$ is the *value function* (i.e., optimal cost-to-go)
- solution \mathbf{u}_x^* is the *optimal control*
- *optimal trajectory* z_x^* dictated by \mathbf{u}_x^*

¹Fleming and Soner. *Controlled Markov Processes and Viscosity Solutions*. 2006.

Pontryagin Maximum Principle (PMP)

Existing Approach

Solve the forward-backward system² for $0 \leq t \leq T$

$$\begin{cases} \partial_t z_x^*(t) = -\nabla_{\mathbf{p}} H(t, z_x^*(t), \mathbf{p}_x(t)), \\ \partial_t \mathbf{p}_x(t) = \nabla_{\mathbf{x}} H(t, z_x^*(t), \mathbf{p}_x(t)), \\ z_x^*(0) = \mathbf{x}, \quad \mathbf{p}_x(T) = \nabla G(z_x^*(T)), \end{cases} \quad (3)$$

where

- Hamiltonian $H(t, \mathbf{x}, \mathbf{p}_x) = \sup_{\mathbf{u}_x \in U} \{-\mathbf{p}_x \cdot f(t, \mathbf{x}, \mathbf{u}_x) - L(t, \mathbf{x}, \mathbf{u}_x)\}$
- adjoint $\mathbf{p}_x: [0, T] \rightarrow \mathbb{R}^d$

then notation-wise, we have $u_x^*(t) = \mathbf{u}^*(t, z_x^*(t), \mathbf{p}_x(t))$

²Pontryagin et al. *The Mathematical Theory of Optimal Processes*. 1962.

Pontryagin Maximum Principle (PMP)

Existing Approach

Solve the forward-backward system² for $0 \leq t \leq T$

$$\begin{cases} \partial_t z_x^*(t) = -\nabla_{\mathbf{p}} H(t, z_x^*(t), \mathbf{p}_x(t)), \\ \partial_t \mathbf{p}_x(t) = \nabla_{\mathbf{x}} H(t, z_x^*(t), \mathbf{p}_x(t)), \\ z_x^*(0) = \mathbf{x}, \quad \mathbf{p}_x(T) = \nabla G(z_x^*(T)), \end{cases} \quad (3)$$

where

- Hamiltonian $H(t, \mathbf{x}, \mathbf{p}_x) = \sup_{\mathbf{u}_x \in U} \{-\mathbf{p}_x \cdot f(t, \mathbf{x}, \mathbf{u}_x) - L(t, \mathbf{x}, \mathbf{u}_x)\}$
- adjoint $\mathbf{p}_x: [0, T] \rightarrow \mathbb{R}^d$

then notation-wise, we have $\mathbf{u}_x^*(t) = \mathbf{u}^*(t, z_x^*(t), \mathbf{p}_x(t))$

Comments

- Local solution method
 - ▶ Solved for a single \mathbf{x}
 - ▶ For a new \mathbf{x} , need to resolve (3)
- Solving the system is difficult and depends on the initial guess $\mathbf{p}_x(0)$ (if using a shooting method)

²Pontryagin et al. *The Mathematical Theory of Optimal Processes*. 1962.

Hamilton-Jacobi-Bellman (HJB)

Existing Approach

Solve the HJB PDE³

(also called *dynamic programming* equations)

$$\begin{cases} -\partial_t \Phi(t, \mathbf{x}) = -H(t, \mathbf{x}, \nabla \Phi(t, \mathbf{x})), \\ \Phi(T, \mathbf{x}) = G(\mathbf{x}) \end{cases} \quad (4)$$

arises from correspondence

$$\mathbf{p}_{\mathbf{x}}(t) = \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}^*(t)) \quad (5)$$

³Bellman. *Dynamic Programming*. 1957.

Hamilton-Jacobi-Bellman (HJB)

Existing Approach

Solve the HJB PDE³

(also called *dynamic programming* equations)

$$\begin{cases} -\partial_t \Phi(t, \mathbf{x}) = -H(t, \mathbf{x}, \nabla \Phi(t, \mathbf{x})), \\ \Phi(T, \mathbf{x}) = G(\mathbf{x}) \end{cases} \quad (4)$$

arises from correspondence

$$\mathbf{p}_{\mathbf{x}}(t) = \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}^*(t)) \quad (5)$$

Comments

- *Global* solution method
 - ▶ Solved for all \mathbf{x}
 - ▶ For a new \mathbf{x} , no recomputation
- Need grids to solve (4), which scale poorly to high-dimensions

³Bellman. *Dynamic Programming*. 1957.

Our Approach

Motivation

Corridor Problem

Want:

- Semi-global solution method (from HJB)
 - ⇒ one model useful for many initial conditions
 - ⇒ method is robust to shocks/disturbances
- High-dimensional (from PMP)
 - ⇒ multi-agent problems provide high dimensionality and are easy to visualize

Semi-Global Solution Method

Robust to Shocks

Want: semi-global Φ (value function)

How to obtain:


- Solve for Hamiltonian H
- Replace adjoint p with $\nabla\Phi$ using (5)
- Use initial states sampled from Gaussian distribution
- Solve

$$\min_{\Phi} \mathbb{E}_{x \sim \mathcal{N}(\mu, \Sigma)} \left\{ \int_0^T L(s, z_x(s), u_x(s)) ds + G(z_x(T)) \right\}$$

s.t.

$$\partial_t z_x(t) = -\nabla_p H(t, z_x(t), \nabla\Phi(t, z_x(t))) = -\nabla\Phi(t, z_x(t))$$

For Corridor




Example:

$$\mu = \begin{bmatrix} -2 \\ -2 \\ 2 \\ -2 \end{bmatrix}, \quad \Sigma = I$$

Penalizers

Empirically Effective in Training



HJt penalizer \Rightarrow few time steps^{4,5}

⁴Yang and Karniadakis. "Potential Flow Generator with L_2 Optimal Transport...". 2020.

⁵Onken et al. "OT-Flow: Fast and Accurate Continuous Normalizing Flows via Optimal Transport". 2020.

Formulation

Rewrite time-integrals as part of the ODE

$$\min_{\Phi} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mu, \Sigma)} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}}, \quad (6)$$

subject to

$$\partial_t \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(t) \\ c_{L,\mathbf{x}}(t) \\ c_{HJt,\mathbf{x}}(t) \end{pmatrix} = \begin{pmatrix} -\nabla_{\mathbf{p}} H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) \\ L_{\mathbf{x}}(t) \\ \left| \partial_t \Phi(t, \mathbf{z}_{\mathbf{x}}(t)) - H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) \right| \end{pmatrix}, \quad \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(0) \\ c_{L,\mathbf{x}}(0) \\ c_{HJt,\mathbf{x}}(0) \end{pmatrix} = \begin{pmatrix} \mathbf{x} \\ 0 \\ 0 \end{pmatrix}.$$

where, by the envelope formula,

$$L_{\mathbf{x}}(t) = \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t)) \cdot \nabla_{\mathbf{p}} H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) - H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t)))$$

Scalars $\beta_1, \beta_2, \beta_3$ are weighted multipliers (NN hyperparameters)

How do we solve this PDE-constrained optimization problem?

How do we solve this PDE-constrained optimization problem?

Blend Neural Networks and Differential Equations

Choose your buzzword: Neural ODEs, Physics-Informed Neural Networks, etc.

Neural Network (NN) Basics

Consider a parameterized function:

$$C = g(z; \theta)$$


where

$z \in \mathbb{R}^d$ is an input item (e.g., the state of the system)

$C \in \mathbb{R}$ is the corresponding output (e.g., the value from Φ)

$\theta \in \mathbb{R}^p$ are the parameters/weights of the model g

Think: Manifold Projection



Single-Layer Example

d - # features

m - width

Features

$$z \in \mathbb{R}^d$$

Weights (θ)

$$K \in \mathbb{R}^{m \times d}$$

$$w \in \mathbb{R}^m$$


$$\text{bias } b \in \mathbb{R}$$

Outputs

$$C \in \mathbb{R}$$

Nonlinearity σ

tanh, sigmoid, etc.



Our Network

A Brief Look Under the Hood

We parameterize the value function

$$\mathbf{a}_0 = \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0),$$

- space-time inputs $\mathbf{s} = (\mathbf{x}, t) \in \mathbb{R}^{d+1}$

⁶He et al. "Deep Residual Learning for Image Recognition". 2016.

Our Network

A Brief Look Under the Hood

We parameterize the value function

where $N(\mathbf{s}) = \mathbf{a}_0 + \sigma(\mathbf{K}_1 \mathbf{a}_0 + \mathbf{b}_1)$,

$$\mathbf{a}_0 = \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0),$$

and

- space-time inputs $\mathbf{s} = (\mathbf{x}, t) \in \mathbb{R}^{d+1}$
- $N(\mathbf{s}) : \mathbb{R}^{d+1} \rightarrow \mathbb{R}^m$ is a residual neural network (ResNet)⁶
- element-wise activation function $\sigma(\mathbf{x}) = \log(\exp(\mathbf{x}) + \exp(-\mathbf{x}))$

⁶He et al. "Deep Residual Learning for Image Recognition". 2016.

Our Network

A Brief Look Under the Hood

We parameterize the value function with

$$\Phi(\mathbf{s}; \boldsymbol{\theta}) = \mathbf{w}^\top N(\mathbf{s}) + \frac{1}{2}\mathbf{s}^\top (\mathbf{A}^\top \mathbf{A})\mathbf{s} + \mathbf{b}^\top \mathbf{s} + c, \quad \text{for } \boldsymbol{\theta} = (\mathbf{w}, \mathbf{A}, \mathbf{b}, c, \mathbf{K}_0, \mathbf{K}_1, \mathbf{b}_0, \mathbf{b}_1)$$

where $N(\mathbf{s}) = \mathbf{a}_0 + \sigma(\mathbf{K}_1 \mathbf{a}_0 + \mathbf{b}_1)$,

$$\mathbf{a}_0 = \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0),$$

and

- space-time inputs $\mathbf{s} = (\mathbf{x}, t) \in \mathbb{R}^{d+1}$
- $N(\mathbf{s}) : \mathbb{R}^{d+1} \rightarrow \mathbb{R}^m$ is a residual neural network (ResNet)⁶
- element-wise activation function $\sigma(\mathbf{x}) = \log(\exp(\mathbf{x}) + \exp(-\mathbf{x}))$
- $\boldsymbol{\theta}$ contains the trainable weights: $\mathbf{w} \in \mathbb{R}^m$, $\mathbf{A} \in \mathbb{R}^{10 \times (d+1)}$, $\mathbf{b} \in \mathbb{R}^{d+1}$, $c \in \mathbb{R}$, $\mathbf{K}_0 \in \mathbb{R}^{m \times (d+1)}$, $\mathbf{K}_1 \in \mathbb{R}^{m \times m}$, and $\mathbf{b}_0, \mathbf{b}_1 \in \mathbb{R}^m$.

⁶He et al. "Deep Residual Learning for Image Recognition". 2016.

Differential Equations

Recall: We are solving

$$\min_{\Phi} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mu, \Sigma)} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}},$$

subject to

$$\partial_t \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(t) \\ c_{L,\mathbf{x}}(t) \\ c_{HJt,\mathbf{x}}(t) \end{pmatrix} = \begin{pmatrix} -\nabla_{\mathbf{p}} H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) \\ L_{\mathbf{x}}(t) \\ \left| \partial_t \Phi(t, \mathbf{z}_{\mathbf{x}}(t)) - H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) \right| \end{pmatrix}, \quad \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(0) \\ c_{L,\mathbf{x}}(0) \\ c_{HJt,\mathbf{x}}(0) \end{pmatrix}, = \begin{pmatrix} \mathbf{x} \\ 0 \\ 0 \end{pmatrix}.$$

Differential Equations

Which is the same as training the neural ODE

$$\min_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mu, \boldsymbol{\Sigma})} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}},$$

subject to

$$\partial_t \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(t) \\ c_{L,\mathbf{x}}(t) \\ c_{HJt,\mathbf{x}}(t) \end{pmatrix} = F(t, \mathbf{z}_{\mathbf{x}}(t), \nabla \Phi(t, \mathbf{z}_{\mathbf{x}}(t); \boldsymbol{\theta})), \quad \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(0) \\ c_{L,\mathbf{x}}(0) \\ c_{HJt,\mathbf{x}}(0) \end{pmatrix}, = \begin{pmatrix} \mathbf{x} \\ 0 \\ 0 \end{pmatrix}.$$

Training and Numerics

Solving the Minimization / Training the Neural ODE:

Iterate through

- ① Solve the ODE
- ② Compute the loss function
- ③ Backpropagate
- ④ Update parameters θ

Training and Numerics

Solving the Minimization / Training the Neural ODE:

Iterate through

- ① Solve the ODE
- ② Compute the loss function
- ③ Backpropagate
- ④ Update parameters θ

ODE solver:

Runge-Kutta 4 \Rightarrow efficient and accurate

Discretize-then-Optimize Approach:^{7,8}

First, discretize the ODE at time points, then optimize over that discretization

As opposed to optimize-then-discretize, e.g., solve Karush-Kuhn-Tucker then discretize

⁷Gholaminejad, Keutzer, and Biros. "ANODE: Unconditionally Accurate Memory-Efficient...". 2019.

⁸Onken and Ruthotto. "Discretize-Optimize vs. Optimize-Discretize for Time-Series...". 2020.

Training and Numerics

Solving the Minimization / Training the Neural ODE:

Iterate through

- ① Solve the ODE
- ② Compute the loss function
- ③ Backpropagate
- ④ Update parameters θ

Loss / Objective Function:

$$J(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}}$$

Training and Numerics

Solving the Minimization / Training the Neural ODE:

Iterate through

- ➊ Solve the ODE
- ➋ Compute the loss function
- ➌ Backpropagate
- ➍ Update parameters θ

Compute gradient with respect to parameters (chain rule)

Use automatic differentiation⁹ to compute $\nabla_{\theta} J$

⁹Nocedal and Wright. *Numerical Optimization*. 2006.

Training and Numerics

Solving the Minimization / Training the Neural ODE:

Iterate through

- ➊ Solve the ODE
- ➋ Compute the loss function
- ➌ Backpropagate
- ➍ Update parameters θ

Use ADAM¹⁰

A stochastic subgradient method with momentum

Empirically, ADAM works well in noisy high-dimensional spaces

¹⁰Kingma and Ba. "Adam: A Method for Stochastic Optimization". 2015.

Results

Small Shock

Large Shock

Baseline Corridor

Running Cost: $L(t, \cdot) = E(\cdot) + \alpha_2 Q(\cdot) + \alpha_3 W(\cdot)$
 Terminal Cost: $G(z) = \frac{\alpha_1}{2} \|z - y\|^2$


Direct Transcription Approach via forward Euler

$$\min_{\{u^{(k)}\}} \quad G\left(z^{(n_t)}\right) + h \sum_{k=0}^{n_t-1} L\left(t^{(k)}, z^{(k)}, u^{(k)}\right)$$

$$\text{s.t. } z^{(k+1)} = z^{(k)} + h f\left(t^{(k)}, z^{(k)}, u^{(k)}\right), \\ z^{(0)} = x$$

where $h=T/n_t$. We use $T=1$ and $n_t=50$.

This is a *local* approach, whereas the NN is *global*



Swap Experiments

Two agents swap positions with hard corridor¹¹

Twelve agents swap positions¹¹



¹¹ Mylvaganam, Sassano, and Astolfi. "A Differential Game Approach to Multi-Agent Collision Avoidance". 2017.

Addressing Curse of Dimensionality¹²

Setup:

- Take subproblems of the 12-agent swap experiment (2, 3, 4, 5, and 6 pairs of agents)
- Train the smallest NN we can that achieves a fixed suboptimality (relative to baseline)

The number of parameters grows linearly with problem dimension d



¹²Bellman. *Dynamic Programming*. 1957.

Swarm Trajectory Planning

50 3-dimensional agents with obstacles¹³


¹³Hönig et al. "Trajectory Planning for Quadrotor Swarms". 2018.

Quadcopter Problem

More complicated dynamics¹⁴

Controls: thrust u , torques $\tau_\psi, \tau_\theta, \tau_\varphi$

$$\dot{z} = f(x, u) \implies \begin{cases} \dot{x} = v_x \\ \dot{y} = v_y \\ \dot{z} = v_z \\ \dot{\psi} = v_\psi \\ \dot{\theta} = v_\theta \\ \dot{\varphi} = v_\varphi \\ \dot{v}_x = \frac{u}{m} f_7(\psi, \theta, \varphi) \\ \dot{v}_y = \frac{u}{m} f_8(\psi, \theta, \varphi) \\ \dot{v}_z = \frac{u}{m} f_9(\theta, \varphi) - g \\ \dot{v}_\psi = \tau_\psi \\ \dot{v}_\theta = \tau_\theta \\ \dot{v}_\varphi = \tau_\varphi \end{cases}$$


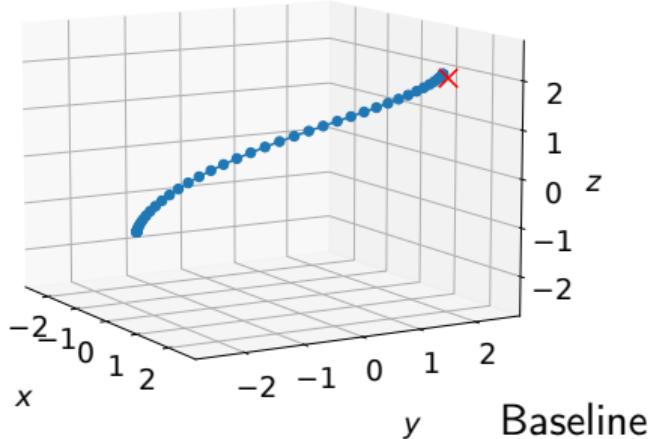



where

$$\begin{cases} f_7(\psi, \theta, \varphi) &= \sin(\psi) \sin(\varphi) + \cos(\psi) \sin(\theta) \cos(\varphi), \\ f_8(\psi, \theta, \varphi) &= -\cos(\psi) \sin(\varphi) + \sin(\psi) \sin(\theta) \cos(\varphi), \\ f_9(\theta, \varphi) &= \cos(\theta) \cos(\varphi). \end{cases}$$

¹⁴Carrillo et al. "Modeling the Quad-Rotor Mini-Rotorcraft". 2013.

Quadcopter Comparison with Baseline



Online/Deployment Timing

Real-time Scenario: at t , want to obtain control to move to $t+1$

Compare

NN: Average cost per Runge-Kutta 4 step ($n_t = 20$ to 50 time steps)

vs.

Baseline: time to obtain 100 gradients for $n_t = 20$ (lower bound for any optimization method)

	Online Time (ms)		Offline (min)
	Baseline lower bound	NN step	NN Train Time
Corridor	2899	4.4	10
Swap 2	2571	4.5	37
Swap 12	1730	3.6	17
Swarm	4026	9.6	57
Quadcopter	3110	5.2	72

Training: on NVIDIA Quadro RTX 8000 GPU.

Online: on 2.6 GHz Intel(R) Xeon(R) CPU E5-4627 core.

Review

- Want to solve
 - ▶ High-Dimensional Control Problems
 - ▶ Semi-Globally
- Combine Pontryagin Maximum Principle and Hamilton-Jacobi-Bellman approaches
- Parameterize the value function Φ with a neural network
- Solve trajectory problem in 150 dimensions
- Solve quadcopter problem with complicated dynamics
- Demonstrate shock-robustness

Conclusions

- Parameterizing Φ
⇒ extrapolation capabilities
- HJB penalizers improve training
- Lagrangian coordinates (no grids) help scalability



DO, L Nurbekyan, X Li, S Wu Fung,
S Osher, L Ruthotto

*A Neural Network Approach Applied to
Multi-Agent Optimal Control*

2021 European Control Conference
arXiv:2011.04757, 2020



DO, L Nurbekyan, X Li, S Wu Fung,
S Osher, L Ruthotto

*A Neural Network Approach for High-Dimensional
Optimal Control*


arXiv:2104.03270, 2021

Code: github.com/donken/NeuralOC

Simulations: imgur.com/a/eWr6sUb

Future Work


- More rigorous experiments with many 12-d quadcopters
- Deployment on actual quadcopters
- Combination with existing methods and sensors



Future Work

- More rigorous experiments with many 12-d quadcopters
- Deployment on actual quadcopters
- Combination with existing methods and sensors

Questions?



References I

- Bellman, Richard (1957). *Dynamic Programming*. Princeton University Press, Princeton, N. J., pp. xxv+342.
- Carrillo, Luis Rodolfo García et al. (2013). "Modeling the Quad-Rotor Mini-Rotorcraft". In: *Quad Rotorcraft Control*. Springer, pp. 23–34.
- Fleming, Wendell H. and H. Mete Soner (2006). *Controlled Markov Processes and Viscosity Solutions*. Second. Vol. 25. Stochastic Modelling and Applied Probability. Springer, New York, pp. xviii+429. ISBN: 978-0387-260457; 0-387-26045-5.
- Gholaminejad, Amir, Kurt Keutzer, and George Biros (2019). "ANODE: Unconditionally Accurate Memory-Efficient Gradients for Neural ODEs". In: *International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 730–736.
- He, Kaiming et al. (2016). "Deep Residual Learning for Image Recognition". In: *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778.
- Hönig, Wolfgang et al. (2018). "Trajectory Planning for Quadrotor Swarms". In: *IEEE Transactions on Robotics* 34.4, pp. 856–869.

References II

- Kingma, Diederik P. and Jimmy Ba (2015). "Adam: A Method for Stochastic Optimization". In: *International Conference on Learning Representations (ICLR)*.
- Mylvgaganam, Thulasi, Mario Sassano, and Alessandro Astolfi (2017). "A Differential Game Approach to Multi-Agent Collision Avoidance". In: *IEEE Transactions on Automatic Control* 62.8, pp. 4229–4235.
- Nocedal, Jorge and Stephen Wright (2006). *Numerical Optimization*. Springer Science & Business Media.
- Onken, Derek and Lars Ruthotto (2020). "Discretize-Optimize vs. Optimize-Discretize for Time-Series Regression and Continuous Normalizing Flows". In: *arXiv:2005.13420*.
- Onken, Derek et al. (2020). "OT-Flow: Fast and Accurate Continuous Normalizing Flows via Optimal Transport". In: *AAAI*.
- Pontryagin, L. S. et al. (1962). *The Mathematical Theory of Optimal Processes*. Translated by K. N. Trirogoff; edited by L. W. Neustadt. Interscience Publishers John Wiley & Sons, Inc. New York-London, pp. viii+360.

References III

Yang, Liu and George Em Karniadakis (2020). "Potential Flow Generator with L_2 Optimal Transport Regularity for Generative Models". In: *IEEE Transactions on Neural Networks and Learning Systems*.