

# A Neural Network Approach for High-Dimensional Optimal Control

Virtual Informal Systems Seminar

February 26, 2021

Derek Onken  
Emory University  
[derekonken.com](http://derekonken.com)



EMORY UNIVERSITY

FOUNDED 1836

# Collaborators and Acknowledgments



Xingjian Li  
Emory



Lars Ruthotto  
Emory



Samy Wu Fung  
UCLA



Levon Nurbekyan  
UCLA



Stan Osher  
UCLA

Funding:



UNITEDHEALTH GROUP®

Special thanks: Organizers and staff of IPAM Long Program MLP 2019 and NVIDIA.

# Overview

- **Background**

- ▶ Problem
- ▶ Pontryagin Maximum Principle (PMP)
- ▶ Hamilton–Jacobi–Bellman Partial Differential Equation (HJB)

- **Mathematical Formulation**

- ▶ Shock-Robustness
- ▶ HJB Penalizers

- **Neural Networks (NNs)**

- ▶ Model Formulation
- ▶ Numerics

- **Results**

- ▶ 150-Dimensional Swarm Trajectory Planning
- ▶ Quadcopter with Complicated Dynamics

- **Conclusion**

# Optimal Control (OC) Problem

## Corridor Problem

Consider two *centrally-controlled* agents that navigate through a corridor/valley between two hills to fixed targets

### Assume

- We have control over the agents' velocities (the *control*)

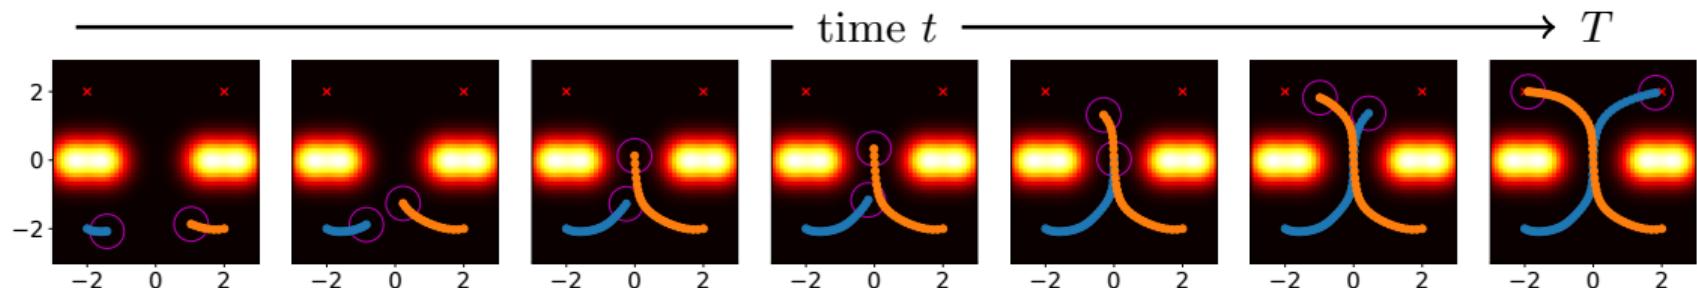
### Want

- Shortest paths, e.g. the geodesics (*optimality*)
- No collisions
- Agents to reach targets at final time

# Multi-Agent Formulation

Consider  $n$  agents initially at  $x_1, \dots, x_n \in \mathbb{R}^q \Rightarrow \mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^d$

Agents follow trajectories  $z_{\mathbf{x}}(t)$  during time  $t \in [0, T]$



Initial

$$z_{\mathbf{x}}(0) = \mathbf{x} = \begin{bmatrix} -2 \\ -2 \\ 2 \\ -2 \end{bmatrix} \left\{ \begin{array}{l} \text{agent 1} \\ \text{agent 2} \end{array} \right.$$

Target

$$\mathbf{y} = \begin{bmatrix} 2 \\ 2 \\ -2 \\ 2 \end{bmatrix}$$

Terminal Cost

$$G(z_{\mathbf{x}}(T)) = \frac{\alpha_1}{2} \|z_{\mathbf{x}}(T) - \mathbf{y}\|^2$$

for multiplier  $\alpha_1 \in \mathbb{R}$

# Trajectories Governed by Differential Equation

The state  $z_x$  depends on the control  $u_x$  and previous state via the system

$$\begin{aligned}\partial_t z_x(t) &= f(t, z_x(t), u_x(t)), \quad z_x(0) = x \\ \text{For Corridor:} \quad &= u_x(t) \quad (\text{the velocity})\end{aligned}\tag{1}$$

where

- time  $t \in [0, T]$
- initial state  $x \in \mathbb{R}^d$
- admissible controls  $U \subset \mathbb{R}^a$
- $f: [0, T] \times \mathbb{R}^d \times U \rightarrow \mathbb{R}^d$  models the evolution of the state  $z_x: [0, T] \rightarrow \mathbb{R}^d$  in response to the control  $u_x: [0, T] \rightarrow U$

## Running Cost

Running costs where  $z_i$  and  $u_i$  are the state and control for the  $i$ th agent, respectively

$$\begin{aligned} L(t, z(t), u(t)) &= E(z(t), u(t)) + \alpha_2 Q(z(t), u(t)) + \alpha_3 W(z(t), u(t)) \\ &= \underbrace{\sum_{i=1}^n E_i(z_i(t), u_i(t))}_{\text{For Corridor: } \frac{1}{2} \|u_i(t)\|^2} + \underbrace{\alpha_2 \sum_{i=1}^n Q_i(z_i(t), u_i(t))}_{\text{sum of Gaussians}} + \underbrace{\alpha_3 \sum_{j \neq i} W_{ij}(z_i(t), z_j(t))}_{\text{piecewise Gaussian repulsion}} \end{aligned}$$

for multipliers  $\alpha_2, \alpha_3 \in \mathbb{R}$  and

- $E_i$  is the energy of an agent,
- $Q_i$  represents any obstacles or terrain,
- $W_{ij}$  are the interaction costs between homogeneous agents  $i$  and  $j$  with radius  $r$

$$W_{ij}(z_i, z_j) = \begin{cases} \exp\left(-\frac{\|z_i - z_j\|_2^2}{2r^2}\right), & \|z_i - z_j\|_2 < 2r \\ 0, & \text{otherwise} \end{cases}$$

## Optimal Control (OC) Problem

$$\begin{aligned} \text{Running Cost: } L(s, \cdot) &= E(\cdot) + \alpha_2 Q(\cdot) + \alpha_3 W(\cdot) \\ \text{Terminal Cost: } G(z_x(T)) &= \frac{\alpha_1}{2} \|z_x(T) - \mathbf{y}\|^2 \end{aligned}$$

**Goal:** Find the control that incurs minimal cost<sup>1</sup>

$$\Phi(t, \mathbf{x}) = \inf_{\mathbf{u}_x} \left\{ \int_t^T L(s, z_x(s), \mathbf{u}_x(s)) ds + G(z_x(T)) \right\} \quad (2)$$

- $\Phi(t, \mathbf{x}) \in \mathbb{R}$  is the *value function* (i.e., optimal cost-to-go)
- solution  $\mathbf{u}_x^*$  is the *optimal control*
- *optimal trajectory*  $z_x^*$  dictated by  $\mathbf{u}_x^*$

<sup>1</sup>Fleming and Soner. *Controlled Markov Processes and Viscosity Solutions*. 2006.

# Pontryagin Maximum Principle (PMP)

## Existing Approach

Solve the forward-backward system<sup>2</sup> for  $0 \leq t \leq T$

$$\begin{cases} \partial_t z_x^*(t) = -\nabla_{\mathbf{p}} H(t, z_x^*(t), \mathbf{p}_x(t)), \\ \partial_t \mathbf{p}_x(t) = \nabla_{\mathbf{x}} H(t, z_x^*(t), \mathbf{p}_x(t)), \\ z_x^*(0) = \mathbf{x}, \quad \mathbf{p}_x(T) = \nabla_{\mathbf{x}} G(z_x^*(T)), \end{cases} \quad (3)$$

where

- Hamiltonian  $H(t, \mathbf{x}, \mathbf{p}_x) = \sup_{\mathbf{u}_x \in U} \{-\mathbf{p}_x \cdot f(t, \mathbf{x}, \mathbf{u}_x) - L(t, \mathbf{x}, \mathbf{u}_x)\}$
- adjoint  $\mathbf{p}_x: [0, T] \rightarrow \mathbb{R}^d$

then notation-wise, we have  $u_x^*(t) = \mathbf{u}^*(t, z_x^*(t), \mathbf{p}_x(t))$

<sup>2</sup>Pontryagin et al. *The Mathematical Theory of Optimal Processes*. 1962.

# Pontryagin Maximum Principle (PMP)

## Existing Approach

Solve the forward-backward system<sup>2</sup> for  $0 \leq t \leq T$

$$\begin{cases} \partial_t z_x^*(t) = -\nabla_p H(t, z_x^*(t), p_x(t)), \\ \partial_t p_x(t) = \nabla_x H(t, z_x^*(t), p_x(t)), \\ z_x^*(0) = x, \quad p_x(T) = \nabla_x G(z_x^*(T)), \end{cases} \quad (3)$$

where

- Hamiltonian  $H(t, x, p_x) = \sup_{u_x \in U} \{-p_x \cdot f(t, x, u_x) - L(t, x, u_x)\}$
- adjoint  $p_x: [0, T] \rightarrow \mathbb{R}^d$

then notation-wise, we have  $u_x^*(t) = u^*(t, z_x^*(t), p_x(t))$

## Comments

- Local solution method
  - ▶ Solved for a single  $x$
  - ▶ For a new  $x$ , need to resolve (3)
- Solving the system is difficult and depends on the initial guess  $p_x(0)$  (if using a shooting method)

<sup>2</sup>Pontryagin et al. *The Mathematical Theory of Optimal Processes*. 1962.

# Hamilton-Jacobi-Bellman (HJB)

## Existing Approach

Solve the HJB PDE<sup>3</sup>

(also called *dynamic programming* equations)

$$\begin{cases} -\partial_t \Phi(t, \mathbf{x}) = -H(t, \mathbf{x}, \nabla_{\mathbf{x}} \Phi(t, \mathbf{x})), \\ \Phi(T, \mathbf{x}) = G(\mathbf{x}) \end{cases} \quad (4)$$

arises from correspondence

$$\mathbf{p}_{\mathbf{x}}(t) = \nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_{\mathbf{x}}^*(t)) \quad (5)$$

<sup>3</sup>Bellman. *Dynamic Programming*. 1957.

# Hamilton-Jacobi-Bellman (HJB)

## Existing Approach

Solve the HJB PDE<sup>3</sup>

(also called *dynamic programming* equations)

$$\begin{cases} -\partial_t \Phi(t, \mathbf{x}) = -H(t, \mathbf{x}, \nabla_{\mathbf{x}} \Phi(t, \mathbf{x})), \\ \Phi(T, \mathbf{x}) = G(\mathbf{x}) \end{cases} \quad (4)$$

arises from correspondence

$$\mathbf{p}_{\mathbf{x}}(t) = \nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_{\mathbf{x}}^*(t)) \quad (5)$$

## Comments

- *Global* solution method
  - ▶ Solved for all  $\mathbf{x}$
  - ▶ For a new  $\mathbf{x}$ , no recomputation
- Need grids to solve (4), which scale poorly to high-dimensions

<sup>3</sup>Bellman. *Dynamic Programming*. 1957.

# Our Approach

Motivation

Corridor Problem

## Want:

- Semi-global solution method (from HJB)
  - ⇒ one model useful for many initial conditions
  - ⇒ method is robust to shocks/disturbances
- High-dimensional (from PMP)
  - ⇒ multi-agent problems provide high dimensionality and are easy to visualize

# Semi-Global Solution Method

Robust to Shocks

**Want:** semi-global  $\Phi$  (value function)

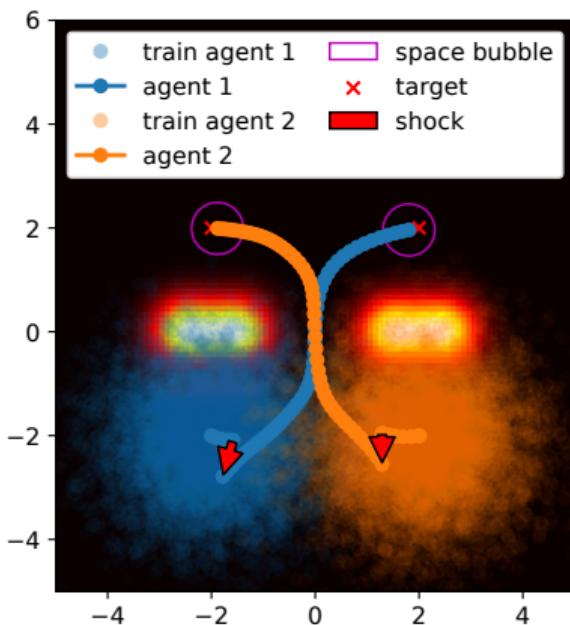
**How to obtain:**

- Solve for Hamiltonian  $H$
- Replace adjoint  $p$  with  $\nabla_x \Phi$  using (5)
- Use initial states sampled from Gaussian distribution
- Solve

$$\min_{\Phi} \mathbb{E}_{\substack{\mathbf{x} \sim \mathcal{N}(\mu, \Sigma)}} \left\{ \int_0^T L(s, \mathbf{z}_x(s), \mathbf{u}_x(s)) ds + G(\mathbf{z}_x(T)) \right\}$$

s.t.

$$\partial_t \mathbf{z}_x(t) = \mathbf{u}_x(t) = -\nabla_x \Phi(t, \mathbf{z}_x(t))$$



Example:

$$\mu = \begin{bmatrix} -2 \\ -2 \\ 2 \\ -2 \end{bmatrix}, \quad \Sigma = I$$

# Penalizers

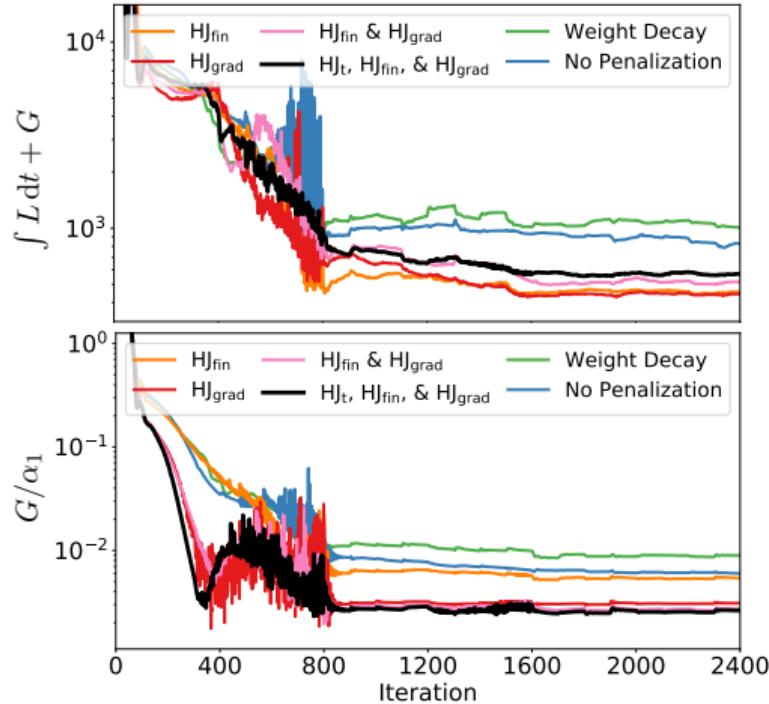
Recall the HJB equations

$$\begin{aligned}-\partial_t \Phi(t, \mathbf{z}_x(t)) &= -H(t, \mathbf{z}_x(t), \nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_x(t))), \\ \Phi(T, \mathbf{z}_x(T)) &= G(\mathbf{z}_x(T))\end{aligned}$$

Make penalizers

$$\begin{aligned}c_{\text{HJt}, \mathbf{x}}(t) &= \int_0^t \left| \partial_s \Phi(s, \mathbf{z}_x(s)) - H(s, \mathbf{z}_x(s), \nabla_{\mathbf{x}} \Phi(s, \mathbf{z}_x(s))) \right| ds \\ c_{\text{HJfin}, \mathbf{x}} &= |\Phi(T, \mathbf{z}_x(T)) - G(\mathbf{z}_x(T))| \\ c_{\text{HJgrad}, \mathbf{x}} &= |\nabla_{\mathbf{x}} \Phi(T, \mathbf{z}_x(T)) - \nabla_{\mathbf{x}} G(\mathbf{z}_x(T))|\end{aligned}$$

## Empirically Effective in Training



HJt penalizer  $\Rightarrow$  few time steps<sup>4,5</sup>

<sup>4</sup>Yang and Karniadakis. "Potential Flow Generator with  $L_2$  Optimal Transport...". 2020.

<sup>5</sup>Onken et al. "OT-Flow: Fast and Accurate Continuous Normalizing Flows via Optimal Transport". 2020.

## Formulation

Rewrite time-integrals as part of the ODE

$$\min_{\Phi} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mu, \Sigma)} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}}, \quad (6)$$

subject to

$$\partial_t \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(t) \\ c_{L,\mathbf{x}}(t) \\ c_{HJt,\mathbf{x}}(t) \end{pmatrix} = \begin{pmatrix} -\nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_{\mathbf{x}}(t)) \\ L(t, \mathbf{z}_{\mathbf{x}}(t), \nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) \\ \left| \partial_t \Phi(t, \mathbf{z}_{\mathbf{x}}(t)) - H(t, \mathbf{z}_{\mathbf{x}}(t), \nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_{\mathbf{x}}(t))) \right| \end{pmatrix}, \quad (7)$$

initialized with  $\mathbf{z}_{\mathbf{x}}(0) = \mathbf{x}$  and  $c_{L,\mathbf{x}}(0) = c_{HJt,\mathbf{x}}(0) = 0$

Scalars  $\beta_1, \beta_2, \beta_3$  are weighted multipliers (NN hyperparameters)

# How do we solve this PDE-constrained optimization problem?

# How do we solve this PDE-constrained optimization problem?

Blend Neural Networks and Differential Equations

Choose your buzzword: Neural ODEs, Physics-Informed Neural Networks, etc.

# Neural Network (NN) Basics

Consider a parameterized function:

$$C = g(z; \theta)$$

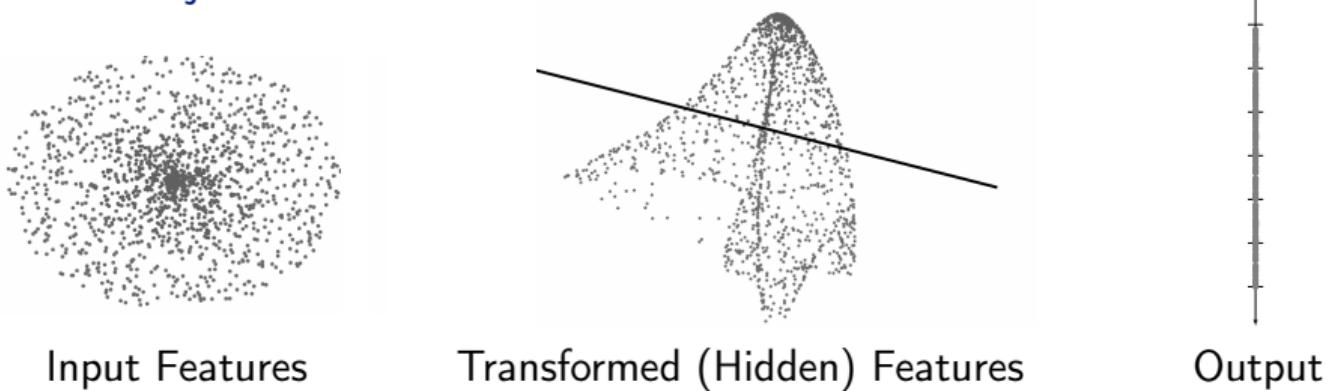
where

$z \in \mathbb{R}^d$  is an input item (e.g., the state of the system)

$C \in \mathbb{R}$  is the corresponding output (e.g., the value from  $\Phi$ )

$\theta \in \mathbb{R}^p$  are the parameters/weights of the model  $g$

Think: Manifold Projection



# Single-Layer Example

$d$  - # features

$m$  - width

Features

$$z \in \mathbb{R}^d$$

Weights ( $\theta$ )

$$K \in \mathbb{R}^{m \times d}$$

$$w \in \mathbb{R}^m$$

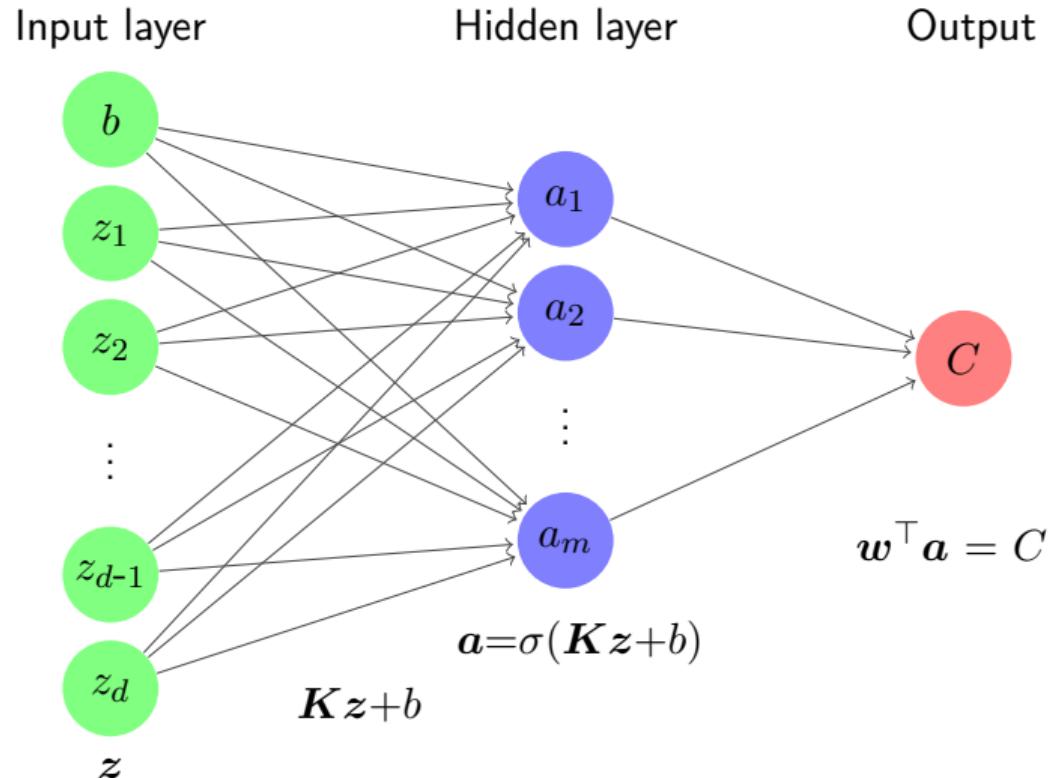
$$\text{bias } b \in \mathbb{R}$$

Outputs

$$C \in \mathbb{R}$$

Nonlinearity  $\sigma$

tanh, sigmoid, etc.



# Our Network

## A Brief Look Under the Hood

We parameterize the value function

$$\mathbf{a}_0 = \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0),$$

- space-time inputs  $\mathbf{s} = (\mathbf{x}, t) \in \mathbb{R}^{d+1}$

<sup>6</sup>He et al. "Deep Residual Learning for Image Recognition". 2016.

# Our Network

## A Brief Look Under the Hood

We parameterize the value function

where  $N(\mathbf{s}) = \mathbf{a}_0 + \sigma(\mathbf{K}_1 \mathbf{a}_0 + \mathbf{b}_1)$ ,

$$\mathbf{a}_0 = \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0),$$

and

- space-time inputs  $\mathbf{s} = (\mathbf{x}, t) \in \mathbb{R}^{d+1}$
- $N(\mathbf{s}) : \mathbb{R}^{d+1} \rightarrow \mathbb{R}^m$  is a residual neural network (ResNet)<sup>6</sup>
- element-wise activation function  $\sigma(\mathbf{x}) = \log(\exp(\mathbf{x}) + \exp(-\mathbf{x}))$

<sup>6</sup>He et al. "Deep Residual Learning for Image Recognition". 2016.

# Our Network

## A Brief Look Under the Hood

We parameterize the value function with

$$\Phi(\mathbf{s}; \boldsymbol{\theta}) = \mathbf{w}^\top N(\mathbf{s}) + \frac{1}{2}\mathbf{s}^\top (\mathbf{A}^\top \mathbf{A})\mathbf{s} + \mathbf{b}^\top \mathbf{s} + c, \quad \text{for } \boldsymbol{\theta} = (\mathbf{w}, \mathbf{A}, \mathbf{b}, c, \mathbf{K}_0, \mathbf{K}_1, \mathbf{b}_0, \mathbf{b}_1)$$

where  $N(\mathbf{s}) = \mathbf{a}_0 + \sigma(\mathbf{K}_1 \mathbf{a}_0 + \mathbf{b}_1)$ ,

$$\mathbf{a}_0 = \sigma(\mathbf{K}_0 \mathbf{s} + \mathbf{b}_0),$$

and

- space-time inputs  $\mathbf{s} = (\mathbf{x}, t) \in \mathbb{R}^{d+1}$
- $N(\mathbf{s}) : \mathbb{R}^{d+1} \rightarrow \mathbb{R}^m$  is a residual neural network (ResNet)<sup>6</sup>
- element-wise activation function  $\sigma(\mathbf{x}) = \log(\exp(\mathbf{x}) + \exp(-\mathbf{x}))$
- $\boldsymbol{\theta}$  contains the trainable weights:  $\mathbf{w} \in \mathbb{R}^m$ ,  $\mathbf{A} \in \mathbb{R}^{10 \times (d+1)}$ ,  $\mathbf{b} \in \mathbb{R}^{d+1}$ ,  $c \in \mathbb{R}$ ,  $\mathbf{K}_0 \in \mathbb{R}^{m \times (d+1)}$ ,  $\mathbf{K}_1 \in \mathbb{R}^{m \times m}$ , and  $\mathbf{b}_0, \mathbf{b}_1 \in \mathbb{R}^m$ .

<sup>6</sup>He et al. "Deep Residual Learning for Image Recognition". 2016.

# Differential Equations

**Recall:** We are solving

$$\min_{\Phi} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mu, \Sigma)} c_{L,\mathbf{x}}(T) + G(z_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}},$$

subject to

$$\partial_t \begin{pmatrix} z_{\mathbf{x}}(t) \\ c_{L,\mathbf{x}}(t) \\ c_{HJt,\mathbf{x}}(t) \end{pmatrix} = \begin{pmatrix} -\nabla_{\mathbf{x}} \Phi(t, z_{\mathbf{x}}(t)) \\ L(t, z_{\mathbf{x}}(t), \nabla_{\mathbf{x}} \Phi(t, z_{\mathbf{x}}(t))) \\ \left| \partial_t \Phi(t, z_{\mathbf{x}}(t)) - H(t, z_{\mathbf{x}}(t), \nabla_{\mathbf{x}} \Phi(t, z_{\mathbf{x}}(t))) \right| \end{pmatrix},$$

initialized with  $z_{\mathbf{x}}(0) = \mathbf{x}$  and  $c_{L,\mathbf{x}}(0) = c_{HJt,\mathbf{x}}(0) = 0$

# Differential Equations

Which is the same as training the neural ODE

$$\min_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\mu, \Sigma)} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}},$$

subject to

$$\partial_t \begin{pmatrix} \mathbf{z}_{\mathbf{x}}(t) \\ c_{L,\mathbf{x}}(t) \\ c_{HJt,\mathbf{x}}(t) \end{pmatrix} = F(t, \mathbf{z}_{\mathbf{x}}(t), \nabla_{\mathbf{x}} \Phi(t, \mathbf{z}_{\mathbf{x}}(t); \boldsymbol{\theta})),$$

initialized with  $\mathbf{z}_{\mathbf{x}}(0) = \mathbf{x}$  and  $c_{L,\mathbf{x}}(0) = c_{HJt,\mathbf{x}}(0) = 0$

# Training and Numerics

## Solving the Minimization / Training the Neural ODE:

Iterate through

- ① Solve the ODE
- ② Compute the loss function
- ③ Backpropagate
- ④ Update parameters  $\theta$

# Training and Numerics

## Solving the Minimization / Training the Neural ODE:

Iterate through

- ① Solve the ODE
- ② Compute the loss function
- ③ Backpropagate
- ④ Update parameters  $\theta$

### ODE solver:

Runge-Kutta 4  $\Rightarrow$  efficient and accurate

### Discretize-then-Optimize Approach:<sup>7,8</sup>

First, discretize the ODE at time points, then optimize over that discretization

As opposed to optimize-then-discretize, e.g., solve Karush-Kuhn-Tucker then discretize

<sup>7</sup>Gholaminejad, Keutzer, and Biros. "ANODE: Unconditionally Accurate Memory-Efficient...". 2019.

<sup>8</sup>Onken and Ruthotto. "Discretize-Optimize vs. Optimize-Discretize for Time-Series...". 2020.

# Training and Numerics

## Solving the Minimization / Training the Neural ODE:

Iterate through

- ① Solve the ODE
- ② Compute the loss function
- ③ Backpropagate
- ④ Update parameters  $\theta$

## Loss / Objective Function:

$$J(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})} c_{L,\mathbf{x}}(T) + G(\mathbf{z}_{\mathbf{x}}(T)) + \beta_1 c_{HJt,\mathbf{x}}(T) + \beta_2 c_{HJfin,\mathbf{x}} + \beta_3 c_{HJgrad,\mathbf{x}}$$

# Training and Numerics

## Solving the Minimization / Training the Neural ODE:

Iterate through

- ➊ Solve the ODE
- ➋ Compute the loss function
- ➌ Backpropagate
- ➍ Update parameters  $\theta$

### Compute gradient with respect to parameters (chain rule)

Use automatic differentiation<sup>9</sup> to compute  $\nabla_{\theta} J$

<sup>9</sup>Nocedal and Wright. *Numerical Optimization*. 2006.

# Training and Numerics

## Solving the Minimization / Training the Neural ODE:

Iterate through

- ➊ Solve the ODE
- ➋ Compute the loss function
- ➌ Backpropagate
- ➍ Update parameters  $\theta$

### Use ADAM<sup>10</sup>

A stochastic subgradient method with momentum

Empirically, ADAM works well in noisy high-dimensional spaces

<sup>10</sup>Kingma and Ba. "Adam: A Method for Stochastic Optimization". 2015.

# Results

Small Shock

Background

Formulation

Neural Networks

Results

Large Shock

Conclusion

VISS 2021

21 / 30

# Baseline Corridor

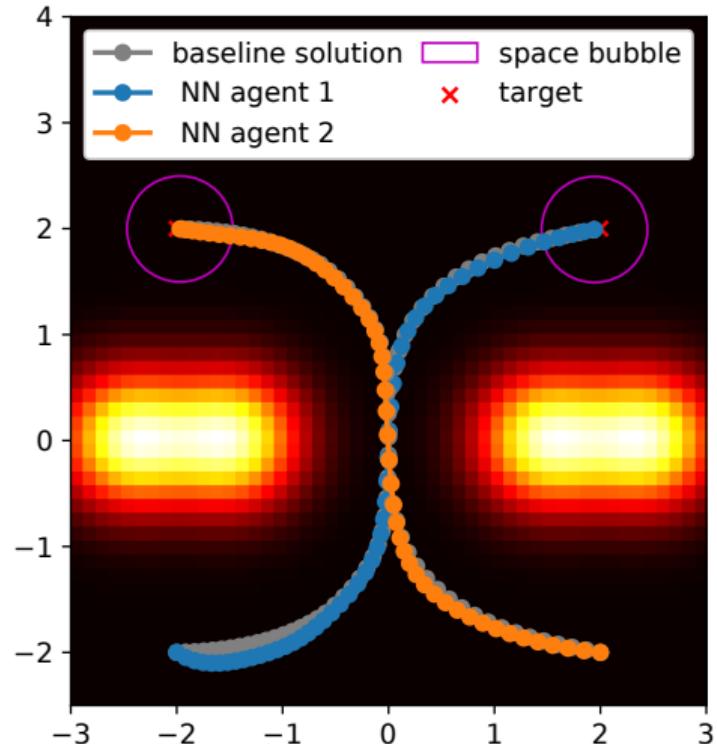
Running Cost:  $L(t, \cdot) = E(\cdot) + \alpha_2 Q(\cdot) + \alpha_3 W(\cdot)$   
 Terminal Cost:  $G(z) = \frac{\alpha_1}{2} \|z - y\|^2$

Discrete optimization approach via forward Euler

$$\begin{aligned} \min_{\{u^{(k)}\}} \quad & G\left(z^{(n_t)}\right) + h \sum_{k=0}^{n_t-1} L\left(t^{(k)}, z^{(k)}, u^{(k)}\right) \\ \text{s.t.} \quad & z^{(k+1)} = z^{(k)} + h f\left(t^{(k)}, z^{(k)}, u^{(k)}\right), \\ & z^{(0)} = x \end{aligned}$$

where  $h=T/n_t$ . We use  $T=1$  and  $n_t=50$ .

This is a *local* approach, whereas the NN is *global*



# Swap Experiments

Two agents swap positions with hard corridor<sup>11</sup>

Twelve agents swap positions<sup>11</sup>

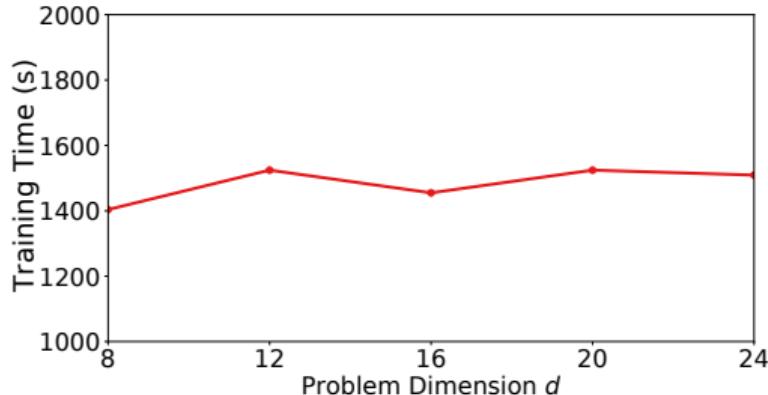
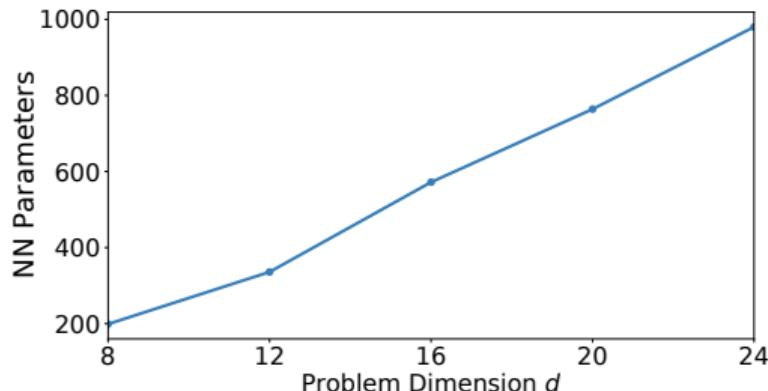
<sup>11</sup> Mylvaganam, Sassano, and Astolfi. "A Differential Game Approach to Multi-Agent Collision Avoidance". 2017.

# Addressing Curse of Dimensionality<sup>12</sup>

## Setup:

- Take subproblems of the 12-agent swap experiment (2, 3, 4, 5, and 6 pairs of agents)
- Train the smallest NN we can that achieves a fixed suboptimality (relative to baseline)

The number of parameters grows linearly with problem dimension  $d$



<sup>12</sup>Bellman. *Dynamic Programming*. 1957.

# Swarm Trajectory Planning

50 3-dimensional agents with obstacles<sup>13</sup>

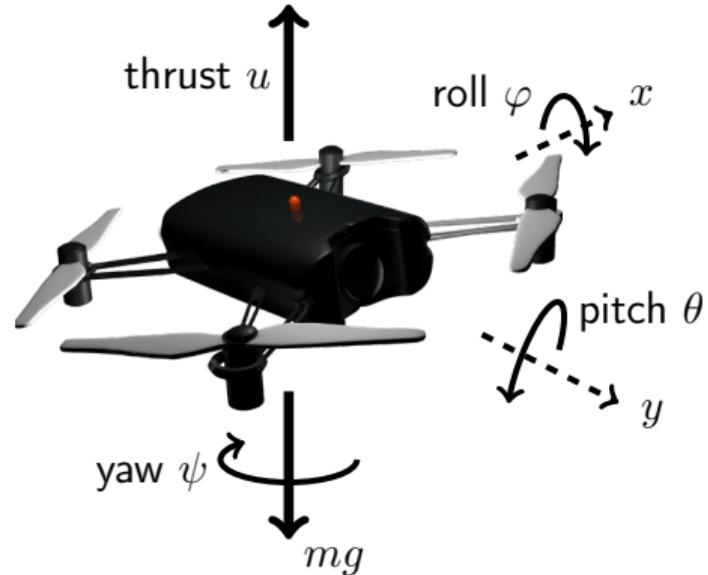
<sup>13</sup>Hönig et al. "Trajectory Planning for Quadrotor Swarms". 2018.

# Quadcopter Problem

More complicated dynamics<sup>14</sup>

Controls: thrust  $u$ , torques  $\tau_\psi, \tau_\theta, \tau_\varphi$

$$\dot{z} = f(x, u) \implies \begin{cases} \dot{x} = v_x \\ \dot{y} = v_y \\ \dot{z} = v_z \\ \dot{\psi} = v_\psi \\ \dot{\theta} = v_\theta \\ \dot{\varphi} = v_\varphi \\ \dot{v}_x = \frac{u}{m} f_7(\psi, \theta, \varphi) \\ \dot{v}_y = \frac{u}{m} f_8(\psi, \theta, \varphi) \\ \dot{v}_z = \frac{u}{m} f_9(\theta, \varphi) - g \\ \dot{v}_\psi = \tau_\psi \\ \dot{v}_\theta = \tau_\theta \\ \dot{v}_\varphi = \tau_\varphi \end{cases}$$

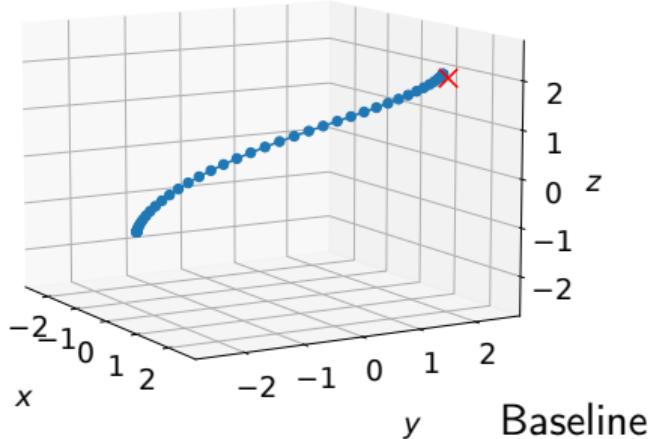
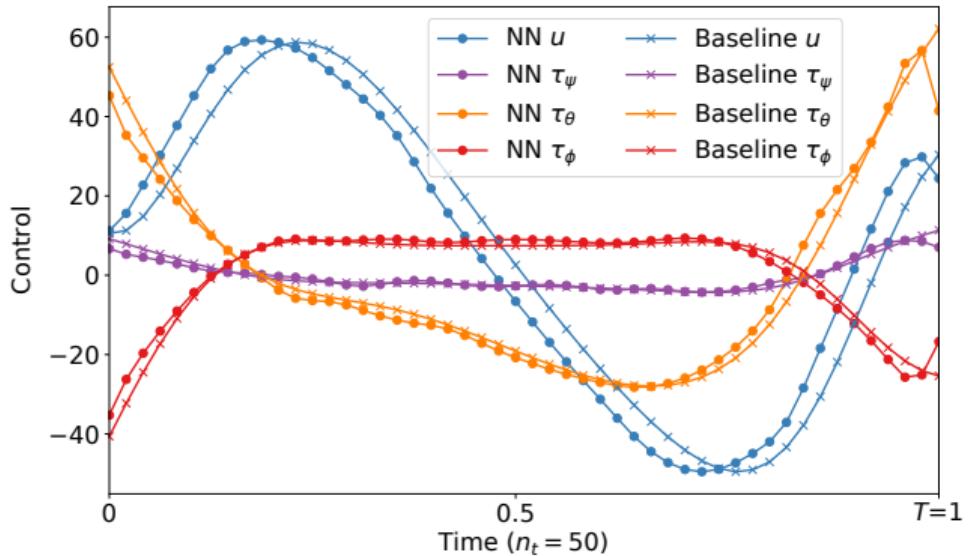


where

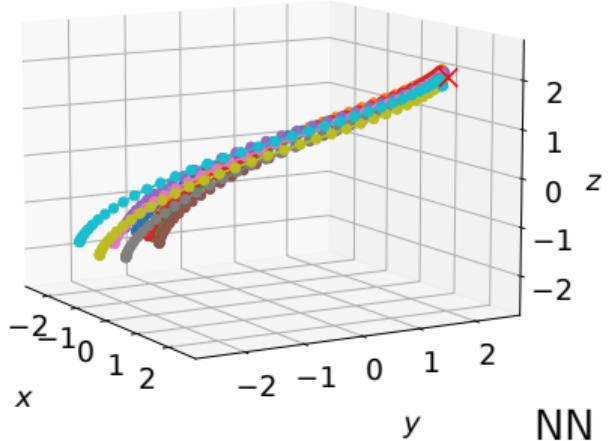
$$\begin{cases} f_7(\psi, \theta, \varphi) &= \sin(\psi) \sin(\varphi) + \cos(\psi) \sin(\theta) \cos(\varphi), \\ f_8(\psi, \theta, \varphi) &= -\cos(\psi) \sin(\varphi) + \sin(\psi) \sin(\theta) \cos(\varphi), \\ f_9(\theta, \varphi) &= \cos(\theta) \cos(\varphi). \end{cases}$$

<sup>14</sup>Carrillo et al. "Modeling the Quad-Rotor Mini-Rotorcraft". 2013.

# Quadcopter Comparison with Baseline



Baseline



NN

# Review

- Want to solve
  - ▶ High-Dimensional Control Problems
  - ▶ Semi-Globally
- Combine Pontryagin Maximum Principle and Hamilton-Jacobi-Bellman approaches
- Parameterize the value function  $\Phi$  with a neural network
- Solve trajectory problem in 150 dimensions
- Solve quadcopter problem with complicated dynamics
- Demonstrate shock-robustness

# Conclusions

- Parameterizing  $\Phi$   
⇒ extrapolation capabilities
- HJB penalizers improve training
- Lagrangian coordinates (no grids) help scalability



DO, L Nurbekyan, X Li, S Wu Fung,  
S Osher, L Ruthotto

*A Neural Network Approach Applied to  
Multi-Agent Optimal Control*

arXiv:2011.04757, 2020

Coming Soon:



DO, L Nurbekyan, X Li, S Wu Fung,  
S Osher, L Ruthotto

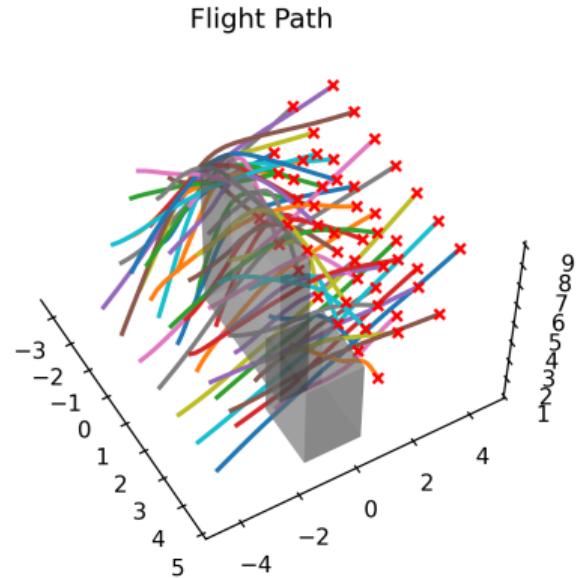
*A Neural Network Approach for High-Dimensional  
Optimal Control*

Code: [github.com/EmoryMLIP/NeuralOC](https://github.com/EmoryMLIP/NeuralOC)

Simulations: [imgur.com/a/eWr6sUb](https://imgur.com/a/eWr6sUb)

# Future Work

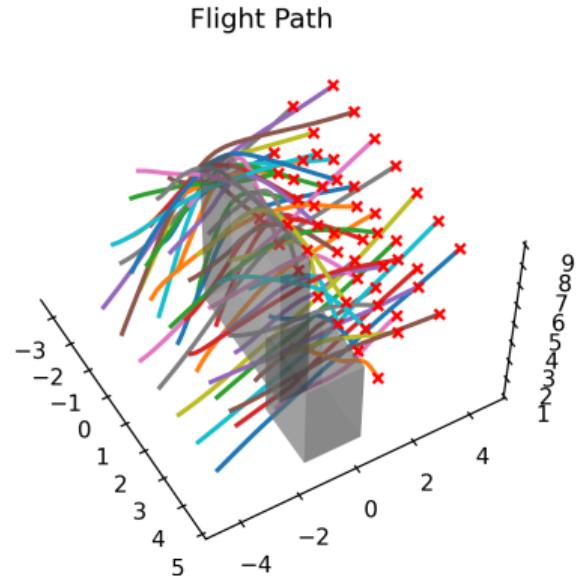
- More rigorous experiments with many 12-d quadcopters
- Deployment on actual quadcopters
- Combination with existing methods and sensors



# Future Work

- More rigorous experiments with many 12-d quadcopters
- Deployment on actual quadcopters
- Combination with existing methods and sensors

Questions?



## References I

- Bellman, Richard (1957). *Dynamic Programming*. Princeton University Press, Princeton, N. J., pp. xxv+342.
- Carrillo, Luis Rodolfo García et al. (2013). “Modeling the Quad-Rotor Mini-Rotorcraft”. In: *Quad Rotorcraft Control*. Springer, pp. 23–34.
- Fleming, Wendell H. and H. Mete Soner (2006). *Controlled Markov Processes and Viscosity Solutions*. Second. Vol. 25. Stochastic Modelling and Applied Probability. Springer, New York, pp. xviii+429. ISBN: 978-0387-260457; 0-387-26045-5.
- Gholaminejad, Amir, Kurt Keutzer, and George Biros (2019). “ANODE: Unconditionally Accurate Memory-Efficient Gradients for Neural ODEs”. In: *International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 730–736.
- He, Kaiming et al. (2016). “Deep Residual Learning for Image Recognition”. In: *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770–778.
- Hönig, Wolfgang et al. (2018). “Trajectory Planning for Quadrotor Swarms”. In: *IEEE Transactions on Robotics* 34.4, pp. 856–869.

## References II

- Kingma, Diederik P. and Jimmy Ba (2015). "Adam: A Method for Stochastic Optimization". In: *International Conference on Learning Representations (ICLR)*.
- Mylvgaganam, Thulasi, Mario Sassano, and Alessandro Astolfi (2017). "A Differential Game Approach to Multi-Agent Collision Avoidance". In: *IEEE Transactions on Automatic Control* 62.8, pp. 4229–4235.
- Nocedal, Jorge and Stephen Wright (2006). *Numerical Optimization*. Springer Science & Business Media.
- Onken, Derek and Lars Ruthotto (2020). "Discretize-Optimize vs. Optimize-Discretize for Time-Series Regression and Continuous Normalizing Flows". In: *arXiv:2005.13420*.
- Onken, Derek et al. (2020). "OT-Flow: Fast and Accurate Continuous Normalizing Flows via Optimal Transport". In: *AAAI*.
- Pontryagin, L. S. et al. (1962). *The Mathematical Theory of Optimal Processes*. Translated by K. N. Trirogoff; edited by L. W. Neustadt. Interscience Publishers John Wiley & Sons, Inc. New York-London, pp. viii+360.

## References III

Yang, Liu and George Em Karniadakis (2020). "Potential Flow Generator with  $L_2$  Optimal Transport Regularity for Generative Models". In: *IEEE Transactions on Neural Networks and Learning Systems*.